作業二

测验, 20 个问题

10 points

1。

Questions 1-2 are about noisy targets.

Consider the bin model for a hypothesis h that makes an error with probability μ in approximating a deterministic target function f (both h and f outputs $\{-1, +1\}$). If we use the same h to approximate a noisy version of f given by

$$P(\mathbf{x}, y) = P(\mathbf{x})P(y|\mathbf{x})$$

$$P(y|\mathbf{x}) = \begin{cases} \lambda & y = f(\mathbf{x}) \\ 1-\lambda & \text{otherwise} \end{cases}$$

What is the probability of error that h makes in approximating the noisy target y?

- ○ $1 - \lambda$
- ○ μ
- ○ $\lambda(1-\mu) + (1-\lambda)\mu$
- ● $\lambda\mu + (1-\lambda)(1-\mu)$
- ○ none of the other choices

---

10 points

2。

作業二
测验, 20 个问题

Following Question 1, with what value of λ will the performance of h be independent of μ ?

- ( ) 0
- ( ) 1
- ( ) 0 or 1
- (•) 0.5
- ( ) none of the other choices

---

10 points

3。
Questions 3-5 are about generalization error, and getting the feel of the bounds numerically. Please use the simple upper bound $N^{d_{\mathrm{vc}}}$ on the growth function $m_{\mathcal{H}}(N)$, assuming that $N \geq 2$ and $d_{vc} \geq 2$.

For an $\mathcal{H}$ with $d_{\mathrm{vc}} = 10$, if you want 95% confidence that your generalization error is at most 0.05 , what is the closest numerical approximation of the sample size that the VC generalization bound predicts?

- ( ) $420,000$
- ( ) $440,000$
- (•) $460,000$
- ( ) $480,000$
- ( ) $500,000$

---

10 points

4。

There are a number of bounds on the generalization error ϵ , all holding with probability at least $1-\delta$. Fix $d_{\rm vc}=50$ and $\delta=0.05$ and plot these bounds as a function of N . Which bound is the tightest (smallest) for very large N , say $N=10,000$?

Note that Devroye and Parrondo & Van den Broek are implicit bounds in ϵ .

- ○ Original VC bound: $\epsilon \le \sqrt{\frac{8}{N}\ln\frac{4m_{\mathcal{H}}(2N)}{\delta}}$
- ○ Rademacher Penalty Bound: $\epsilon \le \sqrt{\frac{2\ln(2Nm_{\mathcal{H}}(N))}{N}}+\sqrt{\frac{2}{N}\ln\frac{1}{\delta}}+\frac{1}{N}$
- ○ Parrondo and Van den Broek: $\epsilon \le \sqrt{\frac{1}{N}\left(2\epsilon+\ln\frac{6m_{\mathcal{H}}(2N)}{\delta}\right)}$
- ● Devroye: $\epsilon \le \sqrt{\frac{1}{2N}\left(4\epsilon(1+\epsilon)+\ln\frac{4m_{\mathcal{H}}(N^2)}{\delta}\right)}$
- ○ Variant VC bound: $\epsilon \le \sqrt{\frac{16}{N}\ln\frac{2m_{\mathcal{H}}(N)}{\sqrt{\delta}}}$

---

10 points

5。
Continuing from Question 4, for small N , say $N=5$, which bound is the tightest (smallest)?

- ○ Original VC bound
- ○ Rademacher Penalty Bound
- ● Parrondo and Van den Broek
- ○ Devroye
- ○ Variant VC bound

作業二
测验, 20 个问题

10 points

6。
In Questions 6-11, you are asked to play with the growth function or VC-dimension of some hypothesis sets.

What is the growth function $m_{\mathcal{H}}(N)$ of "positive-and-negative intervals on $\mathbb{R}$ "? The hypothesis set $\mathcal{H}$ of "positive-and-negative intervals" contains the functions which are $+1$ within an interval $[\ell, r]$ and -1 elsewhere, as well as the functions which are -1 within an interval $[\ell, r]$ and $+1$ elsewhere.

For instance, the hypothesis $h_1(x) = \mathrm{sign}(x(x-4))$ is a negative interval with -1 within $[0, 4]$ and $+1$ elsewhere, and hence belongs to $\mathcal{H}$. The hypothesis $h_2(x) = \mathrm{sign}((x+1)(x)(x-1))$ contains two positive intervals in $[-1, 0]$ and $[1, \infty)$ and hence does not belong to $\mathcal{H}$.

- (●) $N^2 - N + 2$
- ( ) N^2
- ( ) $N^2 + 1$
- ( ) none of the other choices.
- ( ) $N^2 + N + 2$

---

10 points

7。
Continuing from the previous problem, what is the VC-dimension of the hypothesis set of "positive-and-negative intervals on $\mathbb{R}$ "?

- (●) 3
- ( ) 4
- ( ) 5
- ( ) ∞

- 2

10 points

8。
What is the growth function $m_{\mathcal{H}}(N)$ of "positive donuts in $\mathbb{R}^2$ "?

The hypothesis set $\mathcal{H}$ of "positive donuts'' contains hypotheses formed by two concentric circles centered at the origin. In particular, each hypothesis is $+1$ within a "donut" region of $a^2 \le x_1^2 + x_2^2 \le b^2$ and -1 elsewhere. Without loss of generality, we assume $0 < a < b < \infty$.

- $N+1$
- $\binom{N+1}{2}+1$ (selected)
- $\binom{N+1}{3}+1$
- none of the other choices.
- $\binom{N}{2}+1$

10 points

9。
Consider the "polynomial discriminant'' hypothesis set of degree D on $\mathbb{R}$, which is given by

$$\mathcal{H} = \left\{ h_{\mathbf{c}} \;\middle|\; h_{\mathbf{c}}(x) = \text{sign}\left(\sum_{i=0}^{D} c_i x^i \right) \right\}$$

What is the VC-dimension of such an $\mathcal{H}$?

作業二
测验, 20 个问题

- ( ) D
- (•) $D+1$
- ( ) ∞
- ( ) none of the other choices.
- ( ) $D+2$

---

10 points

10。
Consider the "simplified decision trees" hypothesis set on $\mathbb{R}^d$, which is given by

$$\mathcal{H} = \{h_{\mathbf{t},\mathbf{S}} \mid h_{\mathbf{t},\mathbf{S}}(\mathbf{x}) = 2[[\mathbf{v} \in S]] - 1, \text{ where } v_i = [[x_i > t_i]], \\ \mathbf{S} \text{ a collection of vectors in } \{0,1\}^d, \mathbf{t} \in \mathbb{R}^d \quad \}$$

That is, each hypothesis makes a prediction by first using the d thresholds t_i to locate $\mathbf{x}$ to be within one of the 2^d hyper-rectangular regions, and looking up $\mathbf{S}$ to decide whether the region should be $+1$ or -1 .

What is the VC-dimension of the "simplified decision trees" hypothesis set?

- (•) 2^d
- ( ) $2^{d+1}-3$
- ( ) ∞
- ( ) none of the other choices.
- ( ) 2^{d+1}

---

10 points

11。
Consider the "triangle waves'' hypothesis set on $\mathbb{R}$, which is given by

$$\mathcal{H} = \{h_\alpha \mid \; h_\alpha(x) = \text{sign}(|(\alpha x) \bmod 4 - 2| - 1), \alpha \in \mathbb{R}\}$$

Here $(z \bmod 4)$ is a number $z - 4k$ for some integer k such that $z - 4k \in [0, 4)$. For instance, $(11.26 \bmod 4)$ is 3.26, and $(-11.26 \bmod 4)$ is 0.74. What is the VC-dimension of such an $\mathcal{H}$?

- ○ 1
- ○ 2
- ● ∞
- ○ none of the other choices.
- ○ 3

---

10 points

12。
In Questions 12-15, you are asked to verify some properties or bounds on the growth function and VC-dimension.

Which of the following is an upper bounds of the growth function $m_{\mathcal{H}}(N)$ for $N \ge d_{vc} \ge 2$?

- ○ $m_{\mathcal{H}}\left(\lfloor \frac{N}{2} \rfloor\right)$
- ○ $2^{d_{vc}}$
- ● $\min_{1 \le i \le N-1} 2^i m_{\mathcal{H}}(N-i)$
- ○

$\sqrt{N^{d_{vc}}}$

none of the other choices

10 points

13。
Which of the following is not a possible growth functions $m_{\mathcal{H}}(N)$ for some hypothesis set?

- 2^N
- $2^{\lfloor\sqrt{N}\rfloor}$
- 1
- $N^2 - N + 2$
- none of the other choices

10 points

14。
For hypothesis sets $\mathcal{H}_1, \mathcal{H}_2, \ldots, \mathcal{H}_K$ with finite, positive VC-dimensions $d_{vc}(\mathcal{H}_k)$, some of the following bounds are correct and some are not.

Which among the correct ones is the tightest bound on $d_{vc}(\bigcap_{k=1}^{K}\mathcal{H}_k)$, the VC-dimension of the $\textbf{intersection}$ of the sets?

(The VC-dimension of an empty set or a singleton set is taken as zero.)

- $0 \ \le \ d_{vc}(\bigcap_{k=1}^{K}\mathcal{H}_k) \ \le \ \sum_{k=1}^{K} d_{vc}(\mathcal{H}_k)$
- $0 \ \le \ d_{vc}(\bigcap_{k=1}^{K}\mathcal{H}_k) \ \le \ \min\{d_{vc}(\mathcal{H}_k)\}_{k=1}^{K}$
- $0 \ \le \ d_{vc}(\bigcap_{k=1}^{K}\mathcal{H}_k) \ \le \ \max\{d_{vc}(\mathcal{H}_k)\}_{k=1}^{K}$

- ◯ $\min\{d_{vc}(\mathcal{H}_k)\}_{k=1}^{K} \ \leq \ d_{vc}(\bigcap_{k=1}^{K}\mathcal{H}_k) \ \leq \ \max\{d_{vc}(\mathcal{H}_k)\}_{k=1}^{K}$
- ◯ $\min\{d_{vc}(\mathcal{H}_k)\}_{k=1}^{K} \ \leq \ d_{vc}(\bigcap_{k=1}^{K}\mathcal{H}_k) \ \leq \ \sum_{k=1}^{K} d_{vc}(\mathcal{H}_k)$

10
points

15。
For hypothesis sets $\mathcal{H}_1, \mathcal{H}_2, \ldots, \mathcal{H}_K$ with finite, positive VC-dimensions $d_{vc}(\mathcal{H}_k)$, some of the following bounds are correct and some are not.

Which among the correct ones is the tightest bound on $d_{vc}(\bigcup_{k=1}^{K}\mathcal{H}_k)$, the VC-dimension of the **union** of the sets?

- ◯ $0 \ \leq \ d_{vc}(\bigcup_{k=1}^{K}\mathcal{H}_k) \ \leq \ K-1+\sum_{k=1}^{K} d_{vc}(\mathcal{H}_k)$
- ◯ $\min\{d_{vc}(\mathcal{H}_k)\}_{k=1}^{K} \ \leq \ d_{vc}(\bigcup_{k=1}^{K}\mathcal{H}_k) \ \leq \ \sum_{k=1}^{K} d_{vc}(\mathcal{H}_k)$
- ◯ $\max\{d_{vc}(\mathcal{H}_k)\}_{k=1}^{K} \ \leq \ d_{vc}(\bigcup_{k=1}^{K}\mathcal{H}_k) \ \leq \ \sum_{k=1}^{K} d_{vc}(\mathcal{H}_k)$
- ◉ $\max\{d_{vc}(\mathcal{H}_k)\}_{k=1}^{K} \ \leq \ d_{vc}(\bigcup_{k=1}^{K}\mathcal{H}_k) \ \leq \ K-1+\sum_{k=1}^{K} d_{vc}(\mathcal{H}_k)$
- ◯ $0 \ \leq \ d_{vc}(\bigcup_{k=1}^{K}\mathcal{H}_k) \ \leq \ \sum_{k=1}^{K} d_{vc}(\mathcal{H}_k)$

10
points

16。

作業二

测验, 20 个问题

For Questions 16-20, you will play with the decision stump algorithm.

In class, we taught about the learning model of "positive and negative rays'' (which is simply one-dimensional perceptron) for one-dimensional data. The model contains hypotheses of the form:

$$h_{s,\theta}(x) = s \cdot \mathrm{sign}(x - \theta).$$

The model is frequently named the "decision stump'' model and is one of the simplest learning models. As shown in class, for one-dimensional data, the VC dimension of the decision stump model is 2 .

In fact, the decision stump model is one of the few models that we could easily minimize E_{in} efficiently by enumerating all possible thresholds. In particular, for N examples, there are at most $2N$ dichotomies (see page 22 of lecture 5 slides), and thus at most $2N$ different E_{in} values. We can then easily choose the dichotomy that leads to the lowest E_{in} , where ties an be broken by randomly choosing among the lowest E_{in} ones. The chosen dichotomy stands for a combination of some "spot" (range of θ) and s , and commonly the median of the range is chosen as the θ that realizes the dichotomy.

In this problem, you are asked to implement such and algorithm and run your program on an artificial data set. First of all, start by generating a one-dimensional data by the procedure below:

(a) Generate x by a uniform distribution in $[-1, 1]$.

(b) Generate y by $f(x) = \tilde{s}(x)$ + noise where $\tilde{s}(x) = \mathrm{sign}(x)$ and the noise flips the result with 20% probability.

For any decision stump $h_{s,\theta}$ with $\theta \in [-1, 1]$, express $E_{out}(h_{s,\theta})$ as a function of θ and s .

- ○ $0.3 + 0.5s(|\theta| - 1)$
- ○ $0.3 + 0.5s(1 - |\theta|)$
- ● $0.5 + 0.3s(|\theta| - 1)$

- [ ] $0.5 + 0.3s(1 - |\theta|)$
- [ ] none of the other choices

---

10 points

17。
Generate a data set of size 20 by the procedure above and run the one-dimensional decision stump algorithm on the data set. Record E_{in} and compute E_{out} with the formula above. Repeat the experiment (including data generation, running the decision stump algorithm, and computing E_{in} and E_{out}) $5,000$ times. What is the average E_{in} ? Please choose the closest option.

- [ ] 0.05
- [x] 0.15
- [ ] 0.25
- [ ] 0.35
- [ ] 0.45

---

10 points

18。
Continuing from the previous question, what is the average E_{out} ? Please choose the closest option.

- [ ] 0.05
- [ ] 0.15
- [x] 0.25
- [ ] 0.35
- [ ] 0.45

10
points

19。
Decision stumps can also work for multi-dimensional data. In particular, each decision stump now deals with a specific dimension i , as shown below.

$$h_{s,i,\theta}(\mathbf{x}) = s \cdot \text{sign}(x_i - \theta).$$

Implement the following decision stump algorithm for multi-dimensional data:

a) for each dimension $i = 1, 2, \cdots, d$, find the best decision stump $h_{s,i,\theta}$ using the one-dimensional decision stump algorithm that you have just implemented.

b) return the "best of best"' decision stump in terms of E_{in} . If there is a tie , please randomly choose among the lowest-E_{in} ones

The training data $\mathcal{D}_{train}$ is available at:

https://www.csie.ntu.edu.tw/~htlin/mooc/datasets/mlfound_math/hw2_train.dat

The testing data $\mathcal{D}_{test}$ is available at:

https://www.csie.ntu.edu.tw/~htlin/mooc/datasets/mlfound_math/hw2_test.dat

Run the algorithm on the $\mathcal{D}_{train}$. Report the E_{in} of the optimal decision stump returned by your program. Choose the closest option.

- ○ 0.05
- ○ 0.15
- ● 0.25
- ○

- [ ] 0.35
- [ ] 0.45

10 points

20。
Use the returned decision stump to predict the label of each example within $\mathcal{D}_{test}$. Report an estimate of E_{out} by E_{test} . Please choose the closest option.

- [ ] 0.05
- [ ] 0.15
- [ ] 0.25
- [x] 0.35
- [ ] 0.45

- [x] 我（**伟臣 沈**）了解提交不是我自己完成的作业 将永远不会通过此课程或导致我的 Coursera 帐号被关闭。
了解荣誉准则的更多信息

提交测试